

ICT303 – Advanced Machine Learning and Artificial Intelligence

Topic 3: Perceptron and Multilayer Perceptron (MLP)

Hamid Laga
H.Laga@murdoch.edu.au
Office: 245.1.020

How to Get in Touch with the Teaching Team

- Internal and External Students

- Email: H.Laga@murdoch.edu.au.

- Important

- In any communication, please make sure that you
 - Start the subject of your email with ICT303
 - Include your student ID, name, and the lab slot in which you are enrolled.
 - We will do all our best to answer your queries within 24 hrs.

In this Lecture

- From regression to classification

- Concept
- Activation functions

- Multilayer Perceptron (MLP)

- Definition
- Activation functions (again)
- Hidden layers

- Training MLPs

- Forward and backward propagation

- Summary

- Learning objectives

- Understand and describe difference between classification and regression
- Understand the meaning and importance of activation function
- Understand multi-layer perceptrons and describe their advantage and limitations
- Implement linear regression, classification and MLP using Python, NumPy and PyTorch
- Train and apply MLP to practical problems

- Additional readings

- Chapters 4 and 5 of the textbook, available at: <https://d2l.ai/>

Classification Problem

- Let's start with a simple problem
 - We would like to develop a machine learning model that takes as input a greyscale image and returns whether it is of a “cat” or “not a cat”
- This is a binary classification problem
 - Input to the model
 - A greyscale image
 - Output
 - A binary variable with value 1 to indicate a “cat” and 0 to indicate “not a cat”

Classification Problem

- What is classification

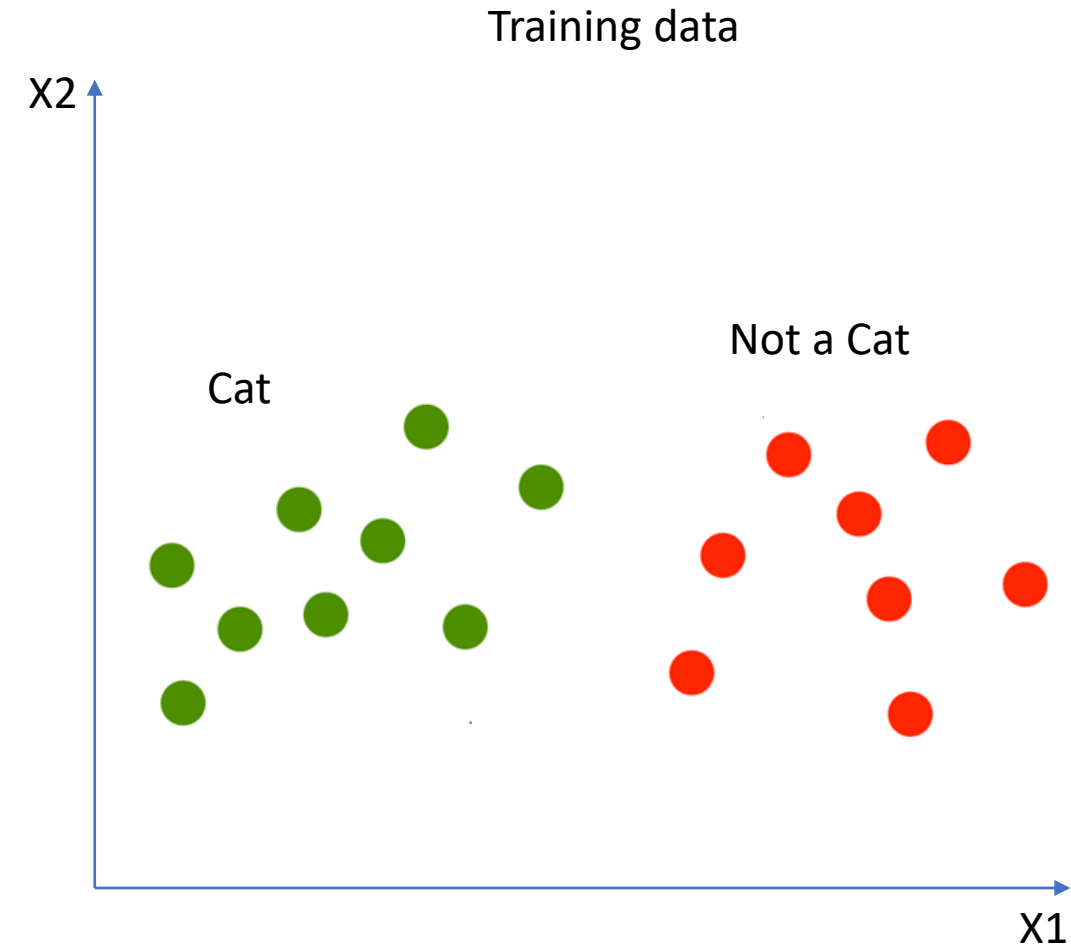
- Given an input (e.g., an image), classify it into one of n categories (e.g., cat, horse, dog, ...)
- Example
 - Object classification from images
 - Face recognition
 - Suspicious behaviour detection
 - ...

- Our objective

- Extend the linear neuron we saw last week to classification problems

Classification Problem


- To start, let's assume that the images are of size 2 x 2
- Input - 2 x 2 grayscale image
 - Each pixel is a grayscale value between x_1, x_2, x_3, x_4
 - Thus, the input is a vector of 4 values
- Output - A binary variable with
 - value 1 to indicate that the image is of a “cat” and
 - Value 0 to indicate that it is “not a cat”



Classification Problem

- To start, let's assume that the images are of size 2 x 2
- Input - 2 x 2 grayscale image
 - Each pixel is a grayscale value between 0 and 1
 - Thus, the input is a vector of 4 values x_1, x_2, x_3, x_4
- Output - A binary variable with
 - value 1 to indicate that the image is of a “cat” and
 - Value 0 to indicate that it is “not a cat”
- Let's use a linear model
 - Similar to linear regression:

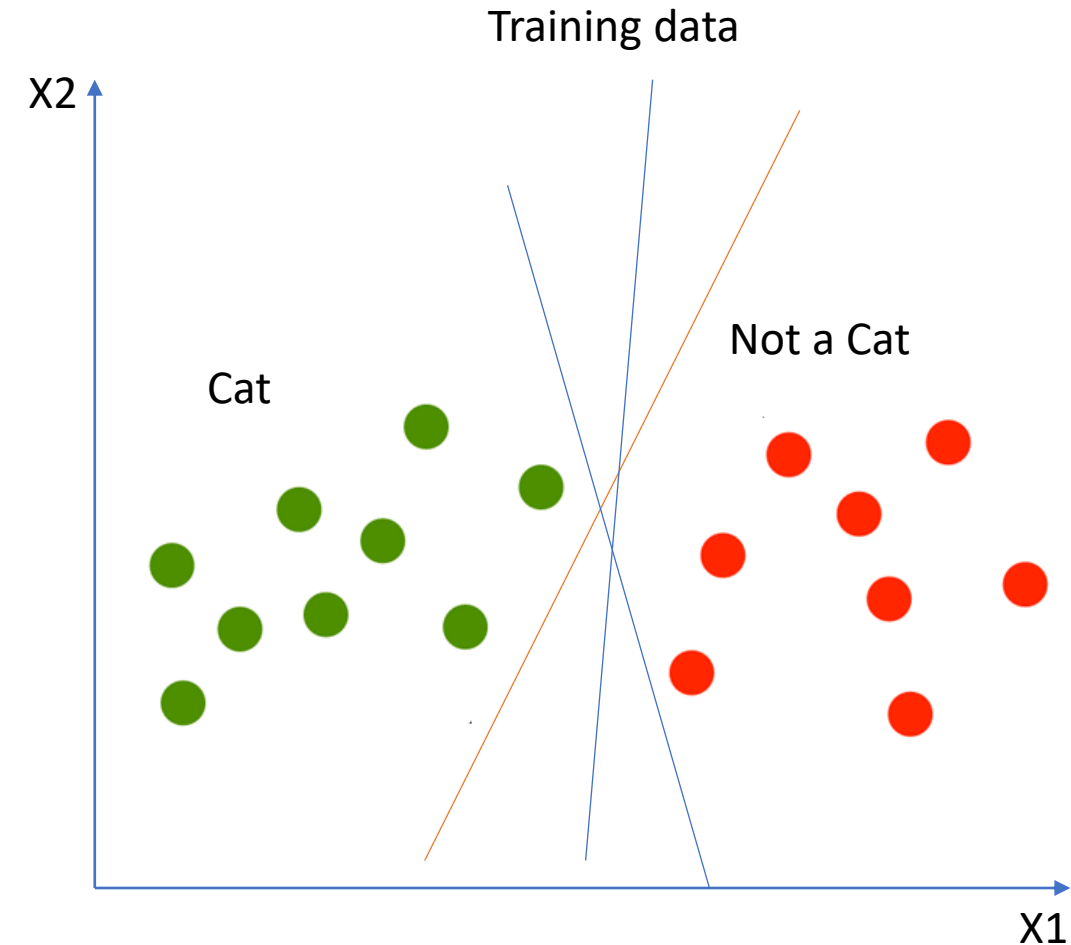
$$O = w_1x_1 + w_2x_2 + w_3x_3 + w_4x_4 + b$$
$$O = \langle \mathbf{W}, \mathbf{x} \rangle + b$$



Classification Problem

- To start, let's assume that the images are of size 2 x 2
- Input - 2 x 2 grayscale image
 - Each pixel is a grayscale value between 0 and 1
 - Thus, the input is a vector of 4 values x_1, x_2, x_3, x_4
- Output - A binary variable with
 - value 1 to indicate that the image is of a “cat” and
 - Value 0 to indicate that it is “not a cat”
- Let's use a linear model
 - Similar to linear regression:

$$O = w_1x_1 + w_2x_2 + w_3x_3 + w_4x_4 + b$$
$$O = \langle \mathbf{W}, \mathbf{x} \rangle + b$$




There are different ways of drawing a boundary between the two classes

Classification Problem

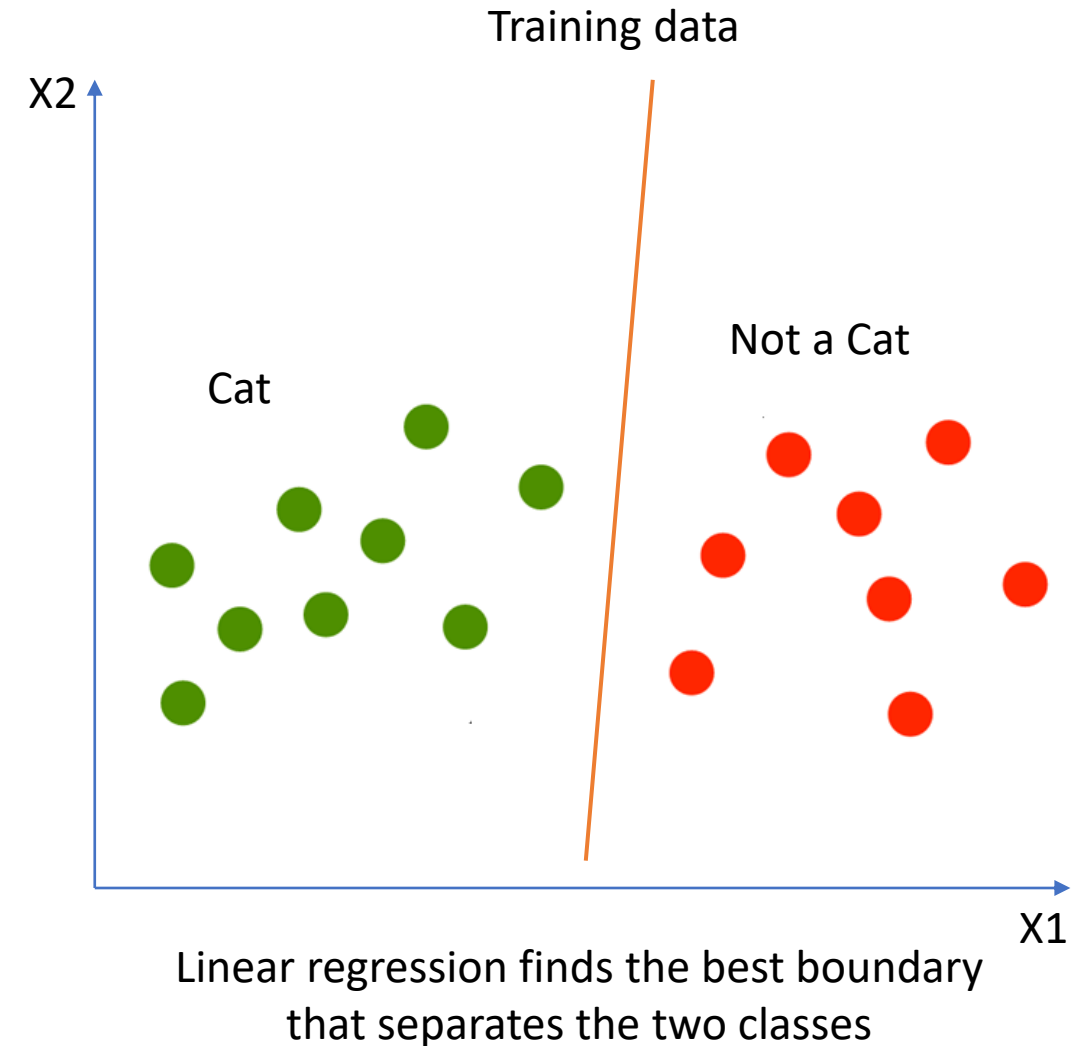
- To start, let's assume that the images are of size 2 x 2
- Input - 2 x 2 grayscale image
 - Each pixel is a grayscale value between 0 and 1
 - Thus, the input is a vector of 4 values x_1, x_2, x_3, x_4
- Output - A binary variable with
 - value 1 to indicate that the image is of a “cat” and
 - Value 0 to indicate that it is “not a cat”
- Let's use a linear model
 - Similar to linear regression:

$$O = w_1x_1 + w_2x_2 + w_3x_3 + w_4x_4 + b$$
$$O = \langle \mathbf{W}, \mathbf{x} \rangle + b$$



Classification Problem


- To start, let's assume that the images are of size 2 x 2
- Input - 2 x 2 grayscale image
 - Each pixel is a grayscale value between 0 and 1
 - Thus, the input is a vector of 4 values x_1, x_2, x_3, x_4
- Output - A binary variable with
 - value 1 to indicate that the image is of a "cat" and
 - Value 0 to indicate that it is "not a cat"
- Let's use a linear model
 - Similar to linear regression:
$$O = w_1x_1 + w_2x_2 + w_3x_3 + w_4x_4 + b$$
$$O = \langle \mathbf{W}, \mathbf{x} \rangle + b$$
 - Linear regression finds the best boundary that separates the classes
 - It does not make a decision



Classification Problem

- To start, let's assume that the images are of size 2 x 2
- Input - 2 x 2 grayscale image
 - Each pixel is a grayscale value between 0 and 1
 - Thus, the input is a vector of 4 values x_1, x_2, x_3, x_4
- Output - A binary variable with
 - value 1 to indicate that the image is of a “cat” and
 - Value 0 to indicate that it is “not a cat”
- Let's use a linear model
 - Similar to linear regression:
$$O = w_1x_1 + w_2x_2 + w_3x_3 + w_4x_4 + b$$
$$O = \langle \mathbf{W}, \mathbf{x} \rangle + b$$
 - Decision:
 - transform the output values so that they are between 0 and 1 using an activation function σ

$$y = \sigma(O)$$



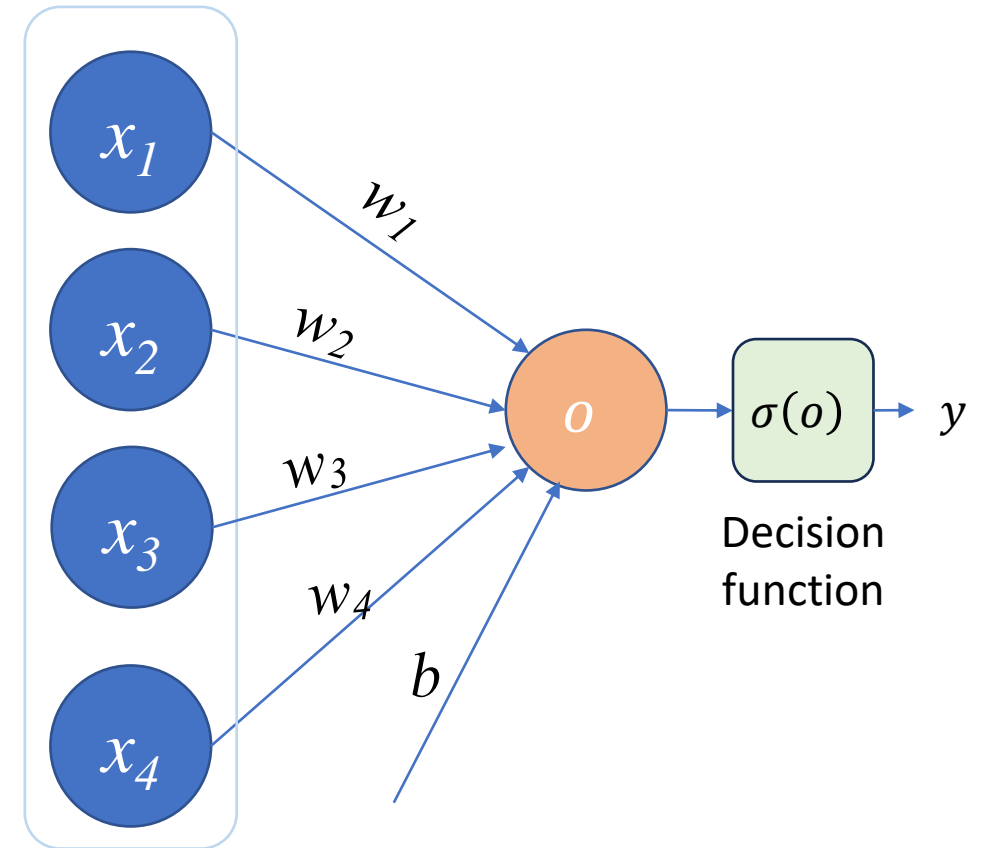
Classification Problem

- To start, let's assume that the images are of size 2 x 2
- Input - 2 x 2 grayscale image
 - Each pixel is a grayscale value between 0 and 1
 - Thus, the input is a vector of 4 values x_1, x_2, x_3, x_4
- Output - A binary variable with
 - value 1 to indicate that the image is of a “cat” and
 - Value 0 to indicate that it is “not a cat”
- Let's use a linear model
 - Similar to linear regression:

$$o = w_1x_1 + w_1x_1 + w_2x_2 + w_3x_3 + w_4x_4 + b$$
$$o = \langle \mathbf{W}, \mathbf{x} \rangle + b$$

- Decision:
 - transform the output values so that they are between 0 and 1 using an activation function σ

$$y = \sigma(o)$$



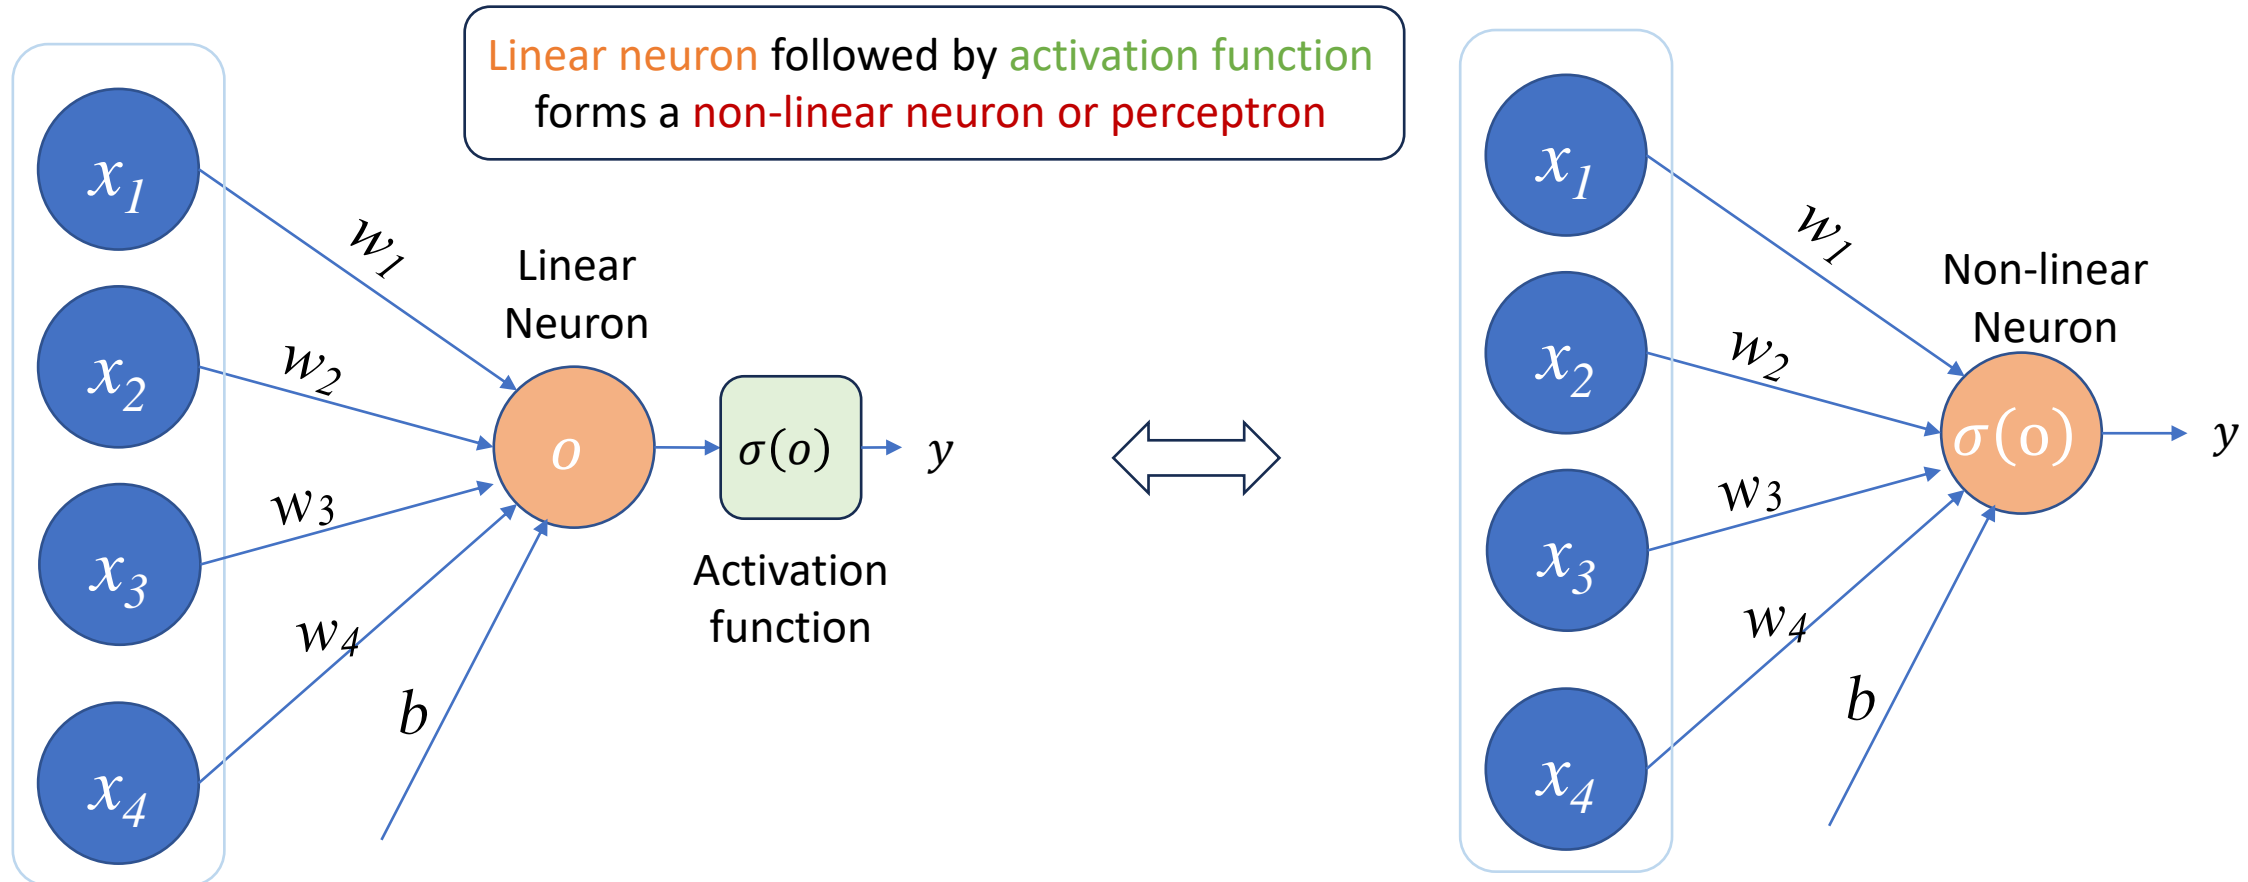
Example of decision, called also activation, functions

$$\sigma(\text{output}) = \begin{cases} 1 & \text{if output} > \epsilon \\ 0 & \text{otherwise} \end{cases}$$

ϵ is a threshold,
e.g., 0.5

Perceptron

- Perceptron is just a linear neuron followed by an activation function
 - It introduces non-linearity – thus sometimes called non-linear neuron



Extension to multiclass classification problem

- We would like to develop a model that takes a greyscale image and returns whether it is of a “cat”, a “car”, or “none of these”

Extension to multiclass classification problem

- We would like to develop a model that takes a greyscale image and returns whether it is of a “cat”, a “car”, or “none of these”

- Input

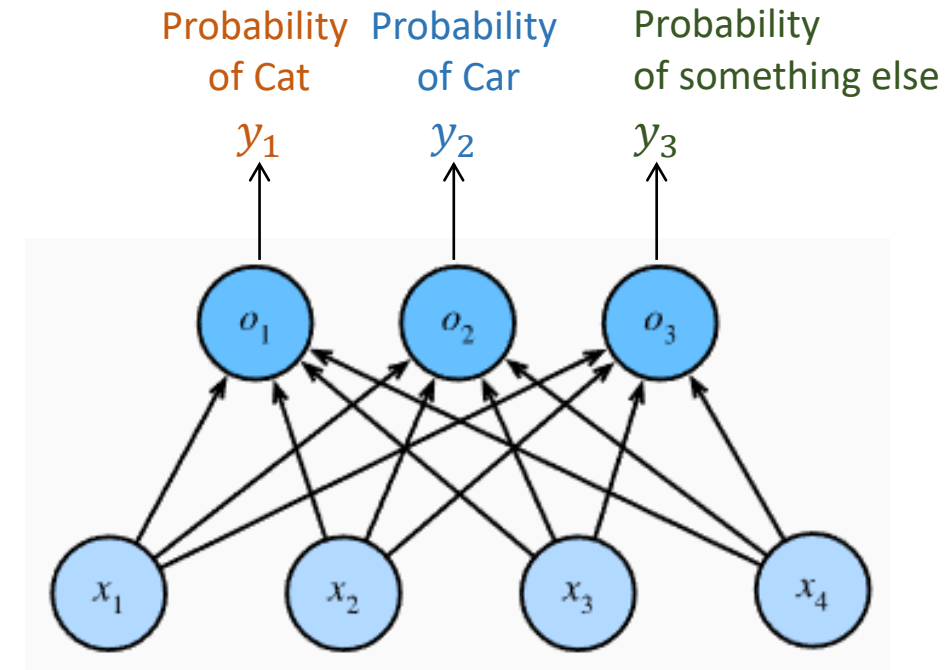
- A greyscale image of size 2 x 2

- Output:

- The probability of being a cat $\rightarrow y_1$
 - The probability of being a car $\rightarrow y_2$
 - The probability of being none of these $\rightarrow y_3$

- Network architecture – two layers

- Input layer has 4 neurons – one for each of the input values
 - The output layer has 3 neurons – one for each of the categories we would like to recognize



Extension to multiclass classification problem

- We would like to develop a model that takes a greyscale image and returns whether it is of a “cat”, a “car”, or “none of these”

- Input


- A greyscale image of size 2 x 2

- Output:

- The probability of being a cat $\rightarrow y_1$
 - The probability of being a car $\rightarrow y_2$
 - The probability of being none of these $\rightarrow y_3$

- Network architecture – two layers

- Input layer has 4 neurons – one for each of the input values
 - The output layer has 3 neurons – one for each of the categories we would like to recognize



Extension to multiclass classification problem

- We would like to develop a model that takes a greyscale image and returns whether it is of a “cat”, a “car”, or “none of these”

- Input

- A greyscale image of size 2 x 2

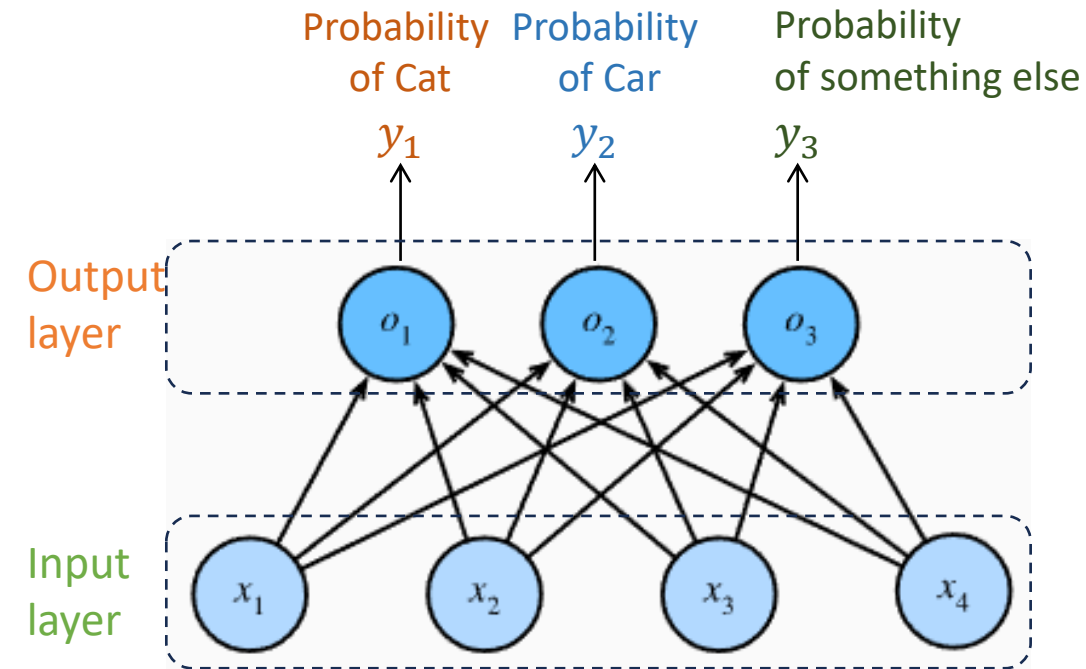
- Output:

- The probability of being a cat $\rightarrow y_1$
 - The probability of being a car $\rightarrow y_2$
 - The probability of being none of these $\rightarrow y_3$

- Network architecture – two layers

- Input layer has 4 neurons – one per input value
 - The output layer has 3 neurons – one per class

$$\begin{aligned}o_1 &= x_1w_{11} + x_2w_{12} + x_3w_{13} + x_4w_{14} + b_1, \\o_2 &= x_1w_{21} + x_2w_{22} + x_3w_{23} + x_4w_{24} + b_2, \\o_3 &= x_1w_{31} + x_2w_{32} + x_3w_{33} + x_4w_{34} + b_3.\end{aligned}$$



Extension to multiclass classification problem

- Simple decision (activation) function

- Can you explain why the simple decision function below won't work for the multiclass problem

$$\sigma(\text{output}) = \begin{cases} 1 & \text{if output} > 0 \\ 0 & \text{otherwise} \end{cases}$$

Extension to multiclass classification problem – Activation Functions

- Max function

- Convert the output to values between 0 and 1
 - $y_1 = \frac{o_1}{o_1 + o_2 + o_3}$
 - $y_2 = \frac{o_2}{o_1 + o_2 + o_3}$
 - $y_3 = \frac{o_3}{o_1 + o_2 + o_3}$
- The final class is the class that has the largest y .

- Softmax function


- Map the output to probabilities by converting the output to values between 0 and 1
 - $y_i = \frac{\exp(o_i)}{\exp(o_1) + \exp(o_2) + \exp(o_3)}$
- Indicates the probability of the input being of that class
- The final class is the class that has the largest y .

- These decision functions are also called activation function

Other Popular Activation Functions

- Sigmoid activation


- A soft version of the previous one
- Maps output to $[0, 1]$



$$\begin{aligned}\sigma(\text{output}) &= \text{sigmoid}(\text{output}) \\ &= \frac{1}{1 + \exp(-\text{output})}\end{aligned}$$

- Tanh activation


- Maps output to $[-1, 1]$



$$\begin{aligned}\sigma(\text{output}) &= \tanh(\text{output}) \\ &= \frac{1 - \exp(-2\text{output})}{1 + \exp(-2\text{output})}\end{aligned}$$

- ReLU activation

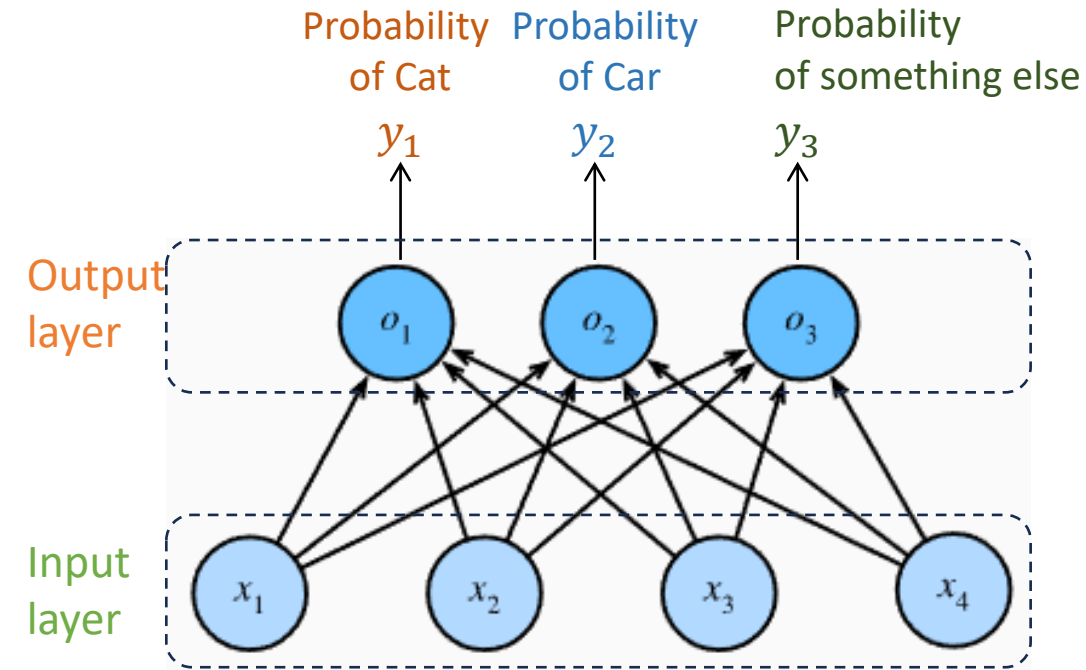
- Zeros negative outputs



$$\begin{aligned}\sigma(\text{output}) &= \text{ReLU}(\text{output}) \\ &= \max(\text{output}, 0)\end{aligned}$$

Multiclass Classification Problem - Summary

- **Input**
 - Images of size $W \times H$
- **Output**
 - $(y_1, y_2, \dots, y_n, y_{n+1})$ where n is the number of classes we would like to recognize plus 1
 - The additional class corresponds to anything else, called background)
- **Model - A network of two layers**
 - Input layer has $W \times H$ neurons
 - Output layer has n neurons
- **Activation function**
 - Softargmax applied to the neurons in the output layer

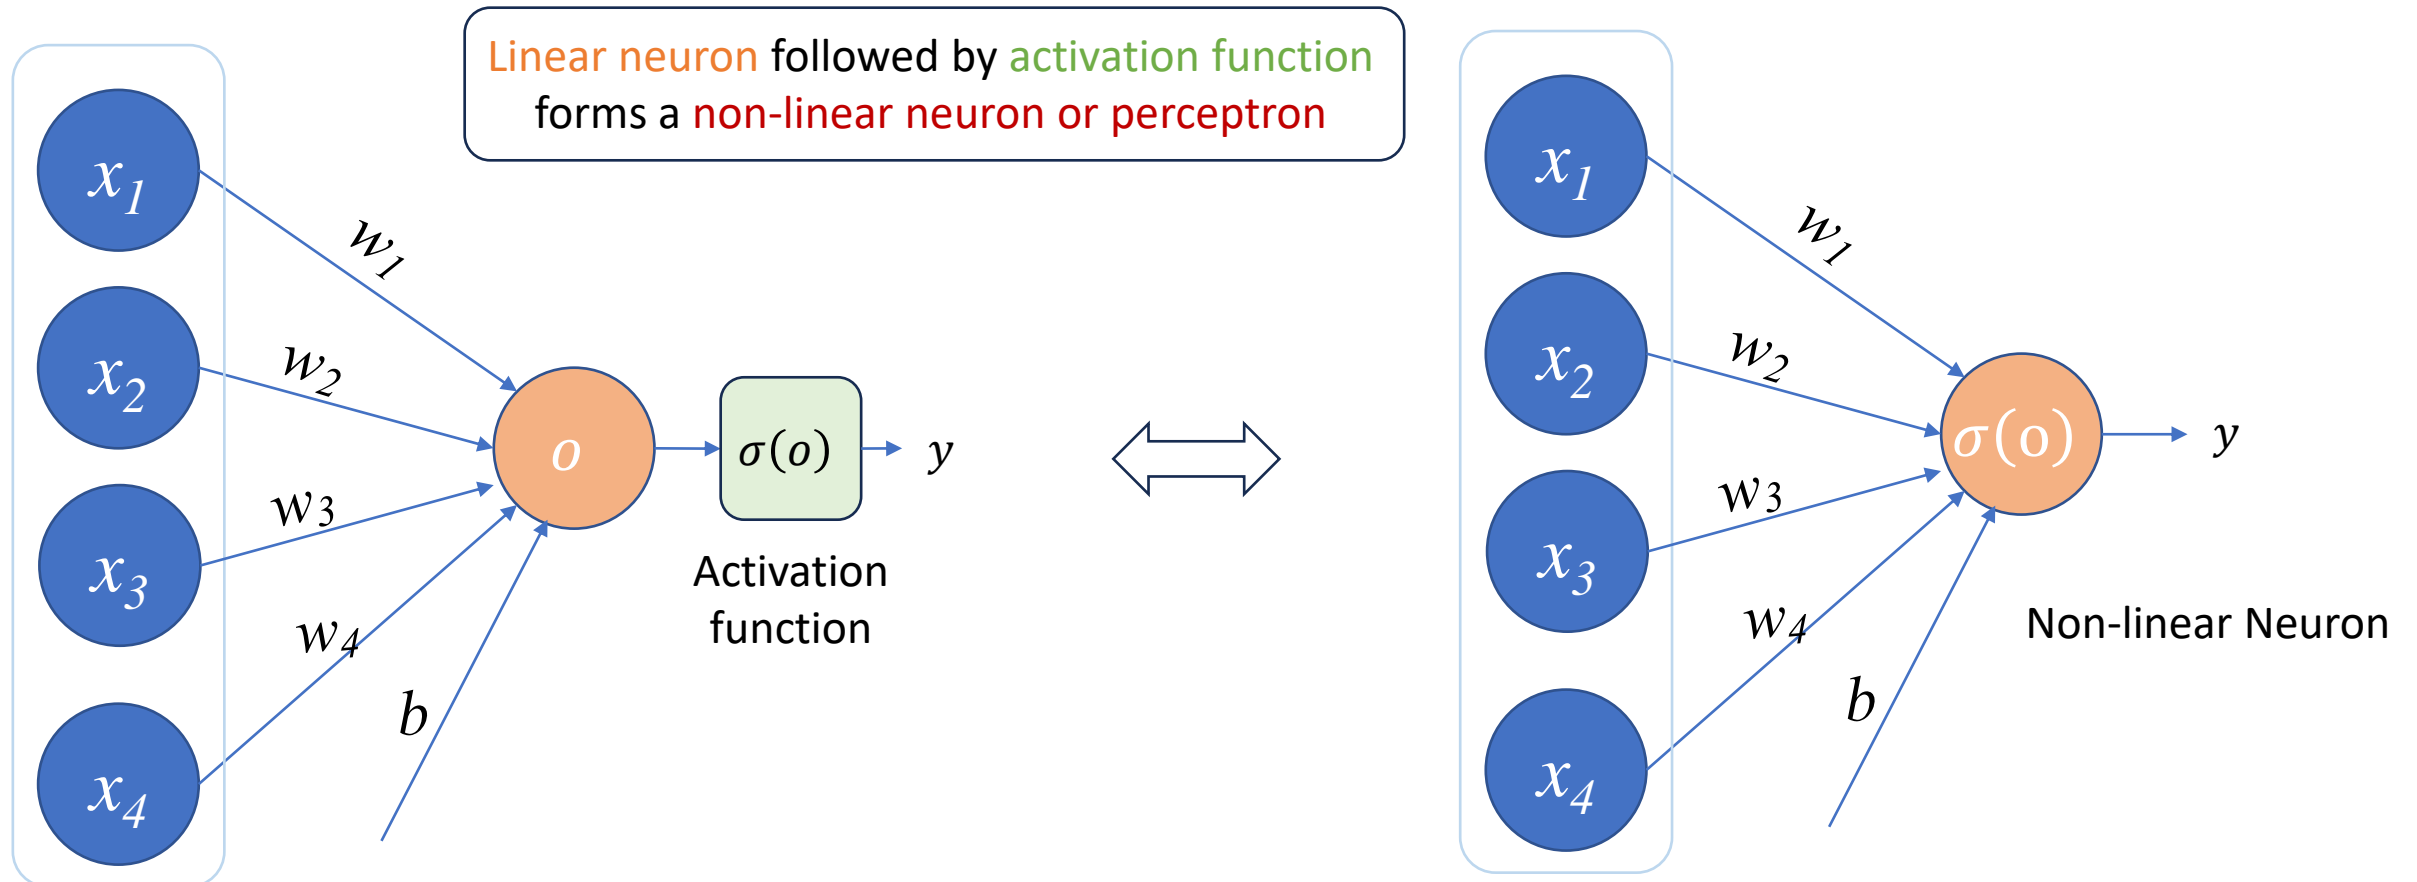


Implementation

- The full implementation will be discussed in the lab
- Questions to ask (and consider)
 - What is the computation time for a layer that has d inputs and q outputs?
 - How many parameters such layer will have?
 - Is it acceptable for real applications, e.g., when dealing with high resolution images and trying to recognize 1000 object categories?

Perceptron

- The activation function introduces non-linearity




Perceptron

- The activation function introduces non-linearity – why is it so important?
 - A single neuron cannot solve some (even simple) problems
 - To solve complex problems, we need to combine multiple neurons
 - Problem
 - Combining multiple linear neurons is the same as having one single linear neuron

A Single Neuron cannot solve complex problems

- XOR Problem (Minsky & Papert, 1969)
 - Input
 - Two binary variables that take values either 0 or 1
 - Out put
 - Their XOR value (see the table on the right)
- This can be modelled as a binary classification problem
 - Class 0 represents the XOR value of 0
 - Class 1 represents the XOR value of 1
- Can you find a line that is able to separate the red class from the green class?

X1	X2	Y
0	0	0
0	1	1
1	0	1
1	1	0



Solution – Use Three Layers of Neurons

- Learning XOR
 - We can write

$$XOR(x_1, x_2) = \underbrace{(x_1 + x_2)}_{\text{OR}} \underbrace{(\overline{x_1 x_2})}_{\text{NAND}}$$

AND

The diagram illustrates the implementation of the XOR function using basic logic gates. The equation $XOR(x_1, x_2) = (x_1 + x_2)(\overline{x_1 x_2})$ is presented. A blue bracket under the term $(x_1 + x_2)$ points down to the label 'OR'. A green bracket under the term $(\overline{x_1 x_2})$ points down to the label 'NAND'. A red bracket over the entire right-hand side of the equation points up to the label 'AND'.

- Show that OR, AND, and NAND (or Not AND) can be implemented using Perceptrons

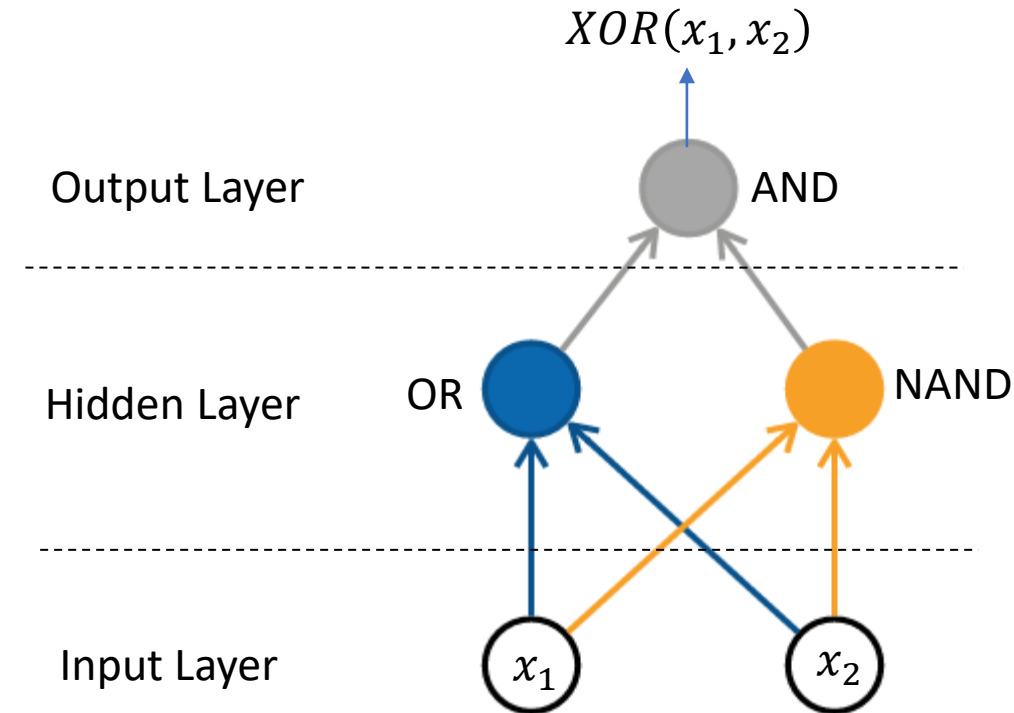
Solution – Use Three Layers of Neurons

- Then XOR can be implemented using multiple layers
 - We can write

$$XOR(x_1, x_2) = \underbrace{(x_1 + x_2)}_{\text{OR}} \underbrace{(\overline{x_1 x_2})}_{\text{NAND}}$$


AND

OR NAND



Multilayer Perceptron – Incorporating Hidden Layers

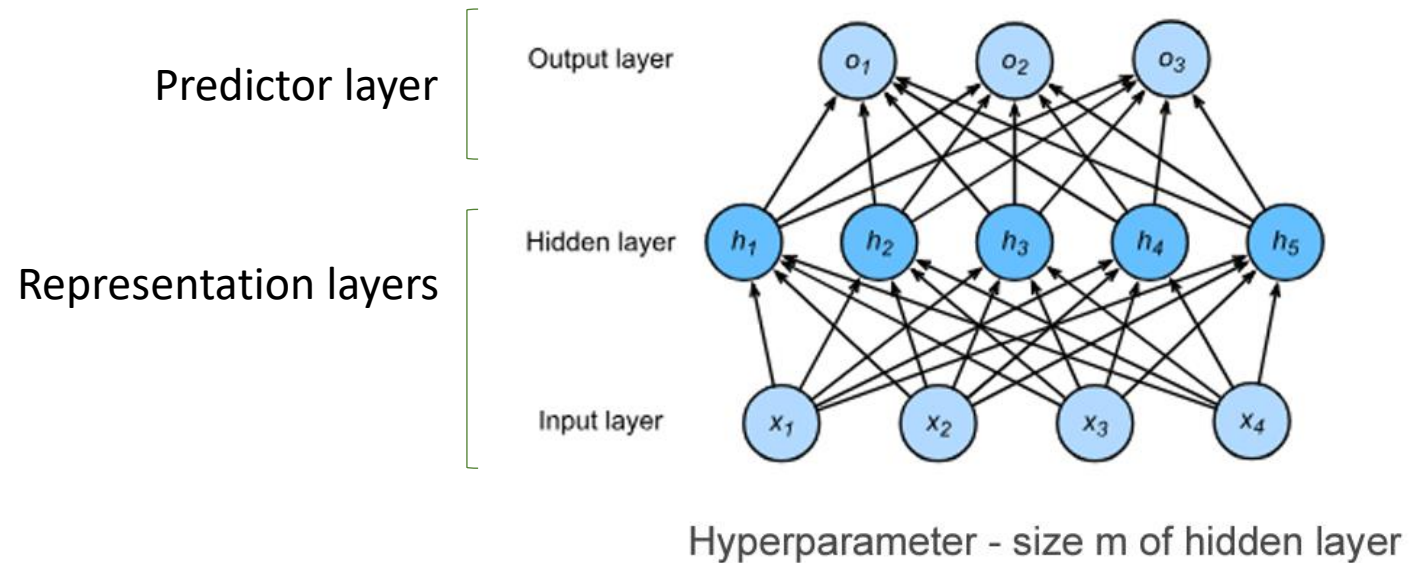
- Solve complex problems by incorporating one or multiple intermediate, called hidden, layers
 - Stack many fully connected layers on top of each other
 - Each layer feeds into the layer above until we generate outputs



Multilayer Perceptron – Incorporating Hidden Layer

- Problem

- Show that cascading multiple **linear** layers is the same as having a network with just one layer of neurons (use a simple example)



Multilayer Perceptron – Incorporating Hidden Layer

- Problem

- Show that cascading multiple **linear** layers is the same as having a network with just one layer of neurons (use a simple example)
- Let h be the output of the hidden layer

$$\mathbf{h} = \mathbf{W}_1 \mathbf{x} + \mathbf{b}_1$$


- Let o be the output of the last layer

$$\mathbf{o} = \mathbf{w}_2^T \mathbf{h} + b_2$$

- Hence:

$$o = \mathbf{w}_2^T \mathbf{W}_1 \mathbf{x} + b'$$

Combining linear layers is
equivalent to on single linear layer



Multilayer Perceptron (MLP) as Universal Approximators


- By applying nonlinear activation function σ to each hidden unit, MLPs become universal approximators

- Input $\mathbf{x} \in \mathbb{R}^n$
- Hidden $\mathbf{W}_1 \in \mathbb{R}^{m \times n}, \mathbf{b}_1 \in \mathbb{R}^m$
- Output $\mathbf{w}_2 \in \mathbb{R}^m, b_2 \in \mathbb{R}$

$$\mathbf{h} = \sigma(\mathbf{W}_1 \mathbf{x} + \mathbf{b}_1)$$

$$\mathbf{o} = \mathbf{w}_2^T \mathbf{h} + b_2$$


σ is an element-wise
activation function



Hyperparameter - size m of hidden layer

Multilayer Perceptron (MLP) as Universal Approximators

- By applying nonlinear activation function σ to each hidden unit, MLPs become universal approximators
- Commonly used activation functions σ
 - In the input layer
 - None
 - In the intermediate (hidden) layer(s)
 - Rectified Linear Unit (ReLU)
 - In the last (output) layer
 - Softmax



Hyperparameter - size m of hidden layer

Multilayer Perceptron (MLP) as Universal Approximators


- By applying nonlinear activation function σ to each hidden unit, MLP become universal approximators

- Input $\mathbf{x} \in \mathbb{R}^n$
- Hidden $\mathbf{W}_1 \in \mathbb{R}^{m \times n}, \mathbf{b}_1 \in \mathbb{R}^m$
- Output $\mathbf{w}_2 \in \mathbb{R}^m, b_2 \in \mathbb{R}$

$$\mathbf{h} = \sigma(\mathbf{W}_1 \mathbf{x} + \mathbf{b}_1)$$

$$\mathbf{o} = \mathbf{w}_2^T \mathbf{h} + b_2$$

$$\mathbf{y} = \text{softmax}(\mathbf{o})$$



Hyperparameter - size m of hidden layer

Multilayer Perceptron (MLP) – Multiple Hidden Layers

- We can cascade multiple hidden layers to approximate complex functions


$$\mathbf{h}_1 = \sigma(\mathbf{W}_1 \mathbf{x} + \mathbf{b}_1)$$

$$\mathbf{h}_2 = \sigma(\mathbf{W}_2 \mathbf{h}_1 + \mathbf{b}_2)$$

$$\mathbf{h}_3 = \sigma(\mathbf{W}_3 \mathbf{h}_2 + \mathbf{b}_3)$$

$$\mathbf{o} = \mathbf{W}_4 \mathbf{h}_3 + \mathbf{b}_4$$

- Hyper parameters
 - Number of hidden layers
 - Number of neurons in each hidden layer



Summary of what we have seen so far

- **Linear neurons**
 - Can be used for linear regression
 - To use them for classification, need to add an activation function (softmax) at the output layer
 - Cascading multiple layers of linear neurons is the same as having one single linear neuron
 - Cannot approximate many functions
- **Multilayer Perceptrons (MLPs) as universal approximators**
 - By adding nonlinear activation functions to the hidden layers, we can approximate any function
 - Can have one or multiple hidden layers (in addition to the input and output layers)
- **Importance of activation functions**
 - Introduce nonlinearity → allows to cascade multiple layers → made NN universal approximators

Training MLPs

- Training is the process of finding the best parameters of the network (i.e., the weights of the different connections)
 - Feed to the network an input, with known (desired) output
 - Compute, in a forward pass, the output of the network
 - Measure the discrepancy between the obtained output and the desired output
 - Update the network weights in such a way the discrepancy is reduced
 - Repeat the process until convergence, i.e., the change in the loss is minimal

Training MLPs

- Training is the process of finding the best parameters of the network (i.e., the weights of the different connections)
 - Feed to the network an input, with known (desired) output → Need training data
 - Compute, in a forward pass, the output of the network
 - Measure the discrepancy between the obtained output and the desired output → Need a loss function
 - Update the network weights in such a way the discrepancy is reduced → Gradient descent algorithm (or any of its latest variants)
 - Repeat the process until convergence, i.e., the change in the loss is minimal → Need to define a stopping criteria

Training – Loss Functions

- Let

- $y = (y_1, y_2, \dots, y_q)$: the ground truth output
- $\hat{y} = (\hat{y}_1, \hat{y}_2, \dots, \hat{y}_q)$: the output produced by the network

- L2 Loss

- $l(y, \hat{y}) = \sqrt{\sum (y_i - \hat{y}_i)^2}$

- L1 Loss

- $l(y, \hat{y}) = \sum |y_i - \hat{y}_i|$

- Cross entropy Loss

- Suitable when using softmax activation

$$l(\mathbf{y}, \hat{\mathbf{y}}) = - \sum_{j=1}^q y_j \log \hat{y}_j.$$

- Its gradient:

$$\partial_{o_j} l(\mathbf{y}, \hat{\mathbf{y}}) = \frac{\exp(o_j)}{\sum_{k=1}^q \exp(o_k)} - y_j = \text{softmax}(\mathbf{o})_j - y_j.$$

Training – Loss Functions

- Use either of the losses defined in the previous slide. Let's call it L
- Update the network weights in such a way the discrepancy (loss) is reduced
- However,
 - We need to ensure that the weights do not grow indefinitely by imposing a constraint s on their norm (this is called L2 regularization)

$$s = \frac{\lambda}{2} \|\mathbf{W}\|^2$$

- Overall loss (called also objective function)

$$J = L + s$$

Training MLPs

- **Forward propagation (or forward pass)**
 - It is the process of propagating the input of the network through all the layers, from the input layer to the output layer
 - It includes the calculation and storage of intermediate variables (including outputs)

- **Once output is obtained for a given training sample,**
 - Compute the objective function
 - Compute its gradient with respect to the parameters
 - Update the parameters

$$h = \rho(z)$$

$$o = W^{(2)}h$$

$$h = \rho(z)$$


$$z = W^{(1)}X$$

X

Output layer

Hidden layer

Input layer



- $W^{(1)}$: weights of the connections from the input layer to the hidden layer
- $W^{(2)}$: weights of the connections from the hidden layer to the output layer

Training MLPs

- **Backpropagation**

- This is the process of computing the gradient of the neural network parameters (i.e., the gradient of the loss function with respect to the NN parameters)
- Introduced in 1980s and was considered as a major breakthrough in neural networks as it allows training, efficiently, neural networks composed of multiple layers

- **When training neural networks**

- Initialize the parameters
- Alternate between forward propagation and backward propagation until convergence
 - Forward propagation sequentially computes and stores intermediate variables (proceeds from the input to the output)
 - Backpropagation sequentially calculates and stores the gradients of intermediate variables and parameters in the reversed order (from output to input)

MLP – Implementation

- Will be covered in the lab

Summary

- **Perceptron**
 - Easy, simple but limited function complexity – cannot represent complex relations between the input and output
- **Multilayer perceptron**
 - Multiple layers add more complexity
 - Nonlinearity is needed → activation function at the hidden layers
- **Activation functions and their importance**
 - Without activation function an MLP is equivalent to a single neuron
- **Training**
 - Loss function
 - Training process

Next Week

- More about training (practical issues to take into account)

Questions
